

E. coli* at Toronto's Beaches: Are Some Beaches Less Safe Than Others?

Marie Curtis Park East and Sunnyside exceed the swimming limit on a third of summer days, more than seven times as often as Gibraltar Point

Sean Murphy

September 26, 2026

Sentence 1: the general area (Toronto's beaches and the daily *E. coli* warning). Sentence 2: data and method (98,370 samples at ten beaches, 2007 to 2026, from Open Data Toronto; days over the limit, a same-day model, yesterday against today). Sentence 3: headline result (Marie Curtis Park East and Sunnyside over the limit on about a third of days, more than seven times Gibraltar Point; yesterday's samples catch 37% of days over the limit). Sentence 4: implication (knowing which beach you are at tells a swimmer more than yesterday's reading).

1 Introduction

[Why *E. coli*] *E. coli* is a sign of faecal pollution. It lives in the gut of people and other warm-blooded animals, so finding it in lake water means that sewage, storm-water runoff or bird droppings have arrived recently, bringing the viruses and bacteria that make swimmers ill [1]. Swimming in water above Toronto's standard raises the risk of ear, eye, nose, throat and skin infections, and swallowed water can cause stomach illness. Young children, older people and people with weakened immune systems are most at risk [2].

[Questions and gap] Placeholder: the two questions (do some beaches go over the limit more often; how well does yesterday's result predict today's) and the gap: no published comparison of the beaches over 20 years, and no check of the one-day delay.

[Findings] Placeholder: 34% and 32% of days at Marie Curtis Park East and Sunnyside against 4% at Gibraltar Point; a warning based on yesterday catches 37% of the days over the limit.

*Code and data are available at: <https://github.com/seandata8/toronto-beaches-ecoli>.

[Why it matters] Placeholder: which beach matters more than which day, and the posting is a weak guide.

[Roadmap] Placeholder: Section 2 describes the data, how a sample becomes a number, and the methods; Section 3 gives the results; Section 4 what they mean and their limits. The appendices cover cleaning and the check on simulated data.

2 Data

2.1 Source and context

[Data source] To reduce water-borne illness among swimmers, Toronto Public Health, Toronto Water, the police Marine Unit and the City’s parks division jointly sample the water at the city’s ten supervised beaches every day from Victoria Day to Labour Day [3]. Samples are taken at four to six fixed sites along each beach (Figure 1), and when the laboratory count of *E. coli* at a beach is too high, Toronto Public Health posts signs warning against swimming [4]. The results are published on the City’s open data portal [3,5] and form the data for this paper.




Figure 1. Each of Toronto’s ten swimming beaches is sampled at four to six sites along the shore. Blue points are the 50 sites this paper uses. The two orange points, sites 60W and GP6, were added in 2026 and sit 14 to 16 km from the beaches they are listed under, so they are left out. Map tiles from Esri.

[Swimming limit] Toronto’s standard of 100 *E. coli* per 100 mL is half that of Ontario and Health Canada: Health Canada raised its guideline from 100 to 200 in 2012, and Ontario follows it, whereas Toronto retained the earlier, stricter value [1,2]. Because 100 is the level at which Toronto posts warnings, it is the limit used throughout this paper.

[Ethics and alternatives] Because the data describe water quality rather than people, they raise no privacy concerns; the ethical stakes lie instead in the advice built on them, since a needless warning costs swimmers a beach day and a missed one exposes them to illness [6]. The City also offers each beach’s daily swimming condition for download [4], but the open data were used here because they give every individual result, allowing each day’s status to be calculated the same way for all twenty years.

[Software] The analysis uses Python [7] with Polars [8], NumPy [9], statsmodels [10], pointblank [11], matplotlib [12] and contextily [13], and the paper is written in Quarto [14].

2.2 Measurement

[Columns] Each row of the dataset records one site on one day: the beach, a site code and its location (Figure 1), the date of collection and the *E. coli* result. To obtain each result, staff collect 200 mL of water about 30 cm below the surface at each of the beach’s sites (five at most beaches, six at Kew Balmy and four at Sunnyside), and a Public Health Ontario laboratory counts the *E. coli* by membrane filtration, in which bacteria trapped on a filter grow overnight into visible colonies that are counted and scaled to a volume of 100 mL [15]. Because the process takes about a day, a posted result describes the previous day’s water [2,4]. When no sample could be taken, usually because of bad weather, the result was left blank [2]; 3.5% of rows are blank.

[Floor and ceiling] The laboratory’s reporting limits compress the data at both ends (Figure 2). Results below 10 are rarely reported, so the 49% of results at exactly 10 are best read as “10 or fewer”. From 2018, high results were increasingly reported as exactly 1,000; earlier results reached 15,820, whereas none exceeded 1,000 in 2026, which suggests that a reporting ceiling was introduced. Because both limits narrow the range, differences among the cleanest beaches and the severity of the worst days are likely understated. The ceiling can also hide a day over the limit: since 2018, 34 days at a beach included a result of 1,000 yet had a daily geometric mean (Section 2.3) below 100, and on about 8 of them a true result of 2,000 or less would have put the day over. The true number of days over the limit since 2018 may therefore be slightly higher than the number calculated from these data.




Figure 2. Every result the City has published, 2007 to 2026. One point is one water sample (98,991 in total), before any cleaning. The vertical axis is logarithmic, so each gridline is 100 times the one below. Points are spread sideways within each year so that they can be seen; the *E. coli* count itself is never moved. Most results are reported in multiples of 10, which is why the low readings fall into separate rows. The laboratory reports no results as zero, and “no colonies detected” is assumed to be reported as 10. From 2018 on, many results are reported as exactly 1,000, and these are assumed to be readings of 1,000 or more. The dashed line marks 100 *E. coli* per 100 mL, the level at which the City posts a swimming warning. The ringed point is the only reading the analysis removes.

[Cleaning] Before analysis, three groups of results were removed: one implausible reading (ringed in Figure 2), the results from two sites that lie far from the beaches they are listed under (Figure 1), and results from dates outside the swimming season. The remaining 98,370 results come from 50 sites on 2,019 days across 20 years; Appendix A gives the details.

[Summary statistics] Most results are low, but two beaches stand apart (Table 1). At eight of the ten beaches the median result (the middle value when all results are ranked) is 10 or 20, and at Gibraltar Point, the cleanest, 61% of results are at the floor of 10 and only 6% exceed 100. At Marie Curtis Park East and Sunnyside, by contrast, the medians are 50 and 60, only 27% and 20% of results are at the floor, and 35% and 34% exceed 100, about six times the share at Gibraltar Point.

Table 1. Results at each beach, 2007 to 2026. A result is the number of *E. coli* per 100 mL in one water sample. Ten is usually the lowest value the laboratory reports, and 100 is Toronto’s limit for swimming. The 90th percentile is the level that one result in ten exceeds. Beach-days are days with at least four results. Beaches run from the largest share of results above 100 to the smallest.

Beach	Sites	Results	Median	90th percentile	At 10	Above 100	Beach-days
Marie Curtis Park East	5	9,612	50	590	27%	35%	1,925
Sunnyside	4	8,006	60	400	20%	34%	1,986
Centre Island	5	9,645	20	220	40%	19%	1,929
Kew Balmy	6	12,047	20	250	47%	19%	2,011
Bluffer’s Park	5	10,006	10	150	55%	13%	2,001
Woodbine	5	10,035	10	110	59%	11%	2,005
Ward’s Island	5	9,757	10	110	58%	10%	1,952
Cherry	5	10,012	10	90	58%	9%	2,003
Hanlan’s Point	5	9,689	10	80	61%	8%	1,939
Gibraltar Point	5	9,561	10	60	61%	6%	1,911

2.3 Methods

[Beach-days] Because Toronto’s limit applies to the geometric mean of a beach’s results on a day rather than to any single result [16], the analysis works with beach-days. The geometric mean of n results is the n th root of their product, or equivalently the average of their logarithms converted back. It suits counts that span orders of magnitude, because a single extreme result does not dominate it: results of 10, 10, 10, 10 and 1,000 give a geometric mean of 25, against an ordinary average of 208. A beach-day was counted only when it had at least four results, rather than five, so that Sunnyside (four sites) is included. This gives 19,662 beach-days, of which 2,816 (14%) were over the limit. A day over the limit is not necessarily a day on which a warning was posted—the data do not record postings. Saleem et al. describe posting decisions based on a geometric mean over two days [6]. In addition, a posted warning reflects the previous day’s water quality, because testing takes about 24 hours (Section 2.2).

[Model] To test whether the differences reflect the beaches rather than the weather on the days each was sampled, each beach was compared only with the other beaches sampled on the same day. For beach b on day d , the model is

$$\log_{10} G_{bd} - \bar{L}_d = \beta_b + \varepsilon_{bd}, \quad (1)$$

where, in Equation 1, G_{bd} is the beach’s geometric mean that day, $\bar{L}_d$ is the average of $\log_{10} G$ over all beaches sampled that day, β_b is the beach’s average gap from the others, and ε_{bd} is

the day-to-day variation the model does not explain. Subtracting each day’s average removes anything that affects all beaches at once, such as a city-wide storm. Because the model works on a log scale, a gap of $\beta_b = 0.3$ means about twice the *E. coli* of the average beach on the same day ($10^{0.3} \approx 2$).

[Model] The standard error of each β_b was calculated from the variation between months rather than between days, because a wet week raises a beach’s levels for several days running. The few sampling days in late May and early September were counted with June and August. Grouping days this way roughly doubles the width of the intervals compared with treating every day as independent. Each estimate is reported with a 95% interval: if the analysis were repeated over many different runs of summers, about 95% of such intervals would contain the true gap. Two beaches were said to differ significantly when a gap as large as the one observed would arise less than 5% of the time if they were in fact alike. Testing all 45 pairs this way would produce about two false differences by chance, so the comparisons were corrected by Holm’s method, which raises the bar each comparison must clear so that the chance of any false difference stays near 5%.

[Consecutive days] Because results take about a day, the sign on display at a beach on any day, called here the posting day, is based on the samples taken the day before. The City posts a “Warning” sign when *E. coli* is over its limit and a “Safe” sign otherwise [2]. To measure how often the sign was right, each posting day was paired with the previous day at the same beach, using only days that were consecutive (18,998 pairs). The sign on the posting day was taken to be “Warning” if the previous day’s geometric mean was over the limit, and “Safe” if it was not. That sign was then compared with the water sampled on the posting day itself.

3 Results

[Beach comparison] Two beaches were over the limit far more often than the other eight (Figure 3). Marie Curtis Park East was over the limit on 34% of its sampled days and Sunnyside on 32%, more than seven times the 4% at Gibraltar Point, the cleanest beach. The other seven beaches ranged from 16% at Kew Balmy to 7% at Cherry and Hanlan’s Point. Sunnyside also stands out in the shape of its distribution: it is the only beach with more days between 30 and 100 than below 30, so its typical day is higher, not only its worst. The ranking is not the result of a few bad years. In 19 of the 20 years, both Marie Curtis and Sunnyside were over the limit more often than every other beach; the exception was 2021, when Kew Balmy (27%) passed Sunnyside (25%). Nor is it an effect of the reporting ceiling introduced in 2018: Marie Curtis was over the limit on 37% of days before 2018 and 31% from 2018, Sunnyside on 32% in both periods, and Gibraltar Point on 5% and 4%.




Figure 3. How often each beach went over the limit, 2007 to 2026. Each bar counts the beach-days whose geometric mean falls in that range (19,662 days in total; days with fewer than four results are left out). Orange bars, right of the dashed line, are days above Toronto's limit of 100 *E. coli* per 100 mL. The lowest bin includes the 17 days whose mean falls below the reporting floor of 10. Panels run from the largest share of days over the limit to the smallest.

[Model estimates] The same-day comparison gives the same ranking (Figure 4). On a typical day, Marie Curtis Park East and Sunnyside each had about 2.2 times the *E. coli* of the average beach sampled that day (95% intervals 2.0 to 2.4 times), whereas Gibraltar Point had about a third less (0.66 times). The two worst beaches did not differ from each other (both 2.19 times), but each differed significantly from every other beach. Of the 45 pairs of beaches, 31 differed after Holm's correction (Appendix C); the 14 that did not were Marie Curtis and Sunnyside, Centre Island and Kew Balmy, and 12 pairs among the six cleanest beaches. Centre Island ranks above Kew Balmy in the model (1.22 against 1.09 times) but below it in the share of days over the limit (14% against 16%), because the model compares typical levels on the same day, whereas the share counts only how often 100 is crossed. The gaps shown are smaller than the true gaps: in simulated data with known differences, the reporting floor and ceiling reduced the estimated gaps to about 65% of their true size, without reversing any of them (Appendix B).




Figure 4. Each beach's *E. coli* on a typical day, as a multiple of the average beach sampled the same day, 2007 to 2026, with 95% intervals. The reporting floor and ceiling make these gaps smaller than the true ones.

[Yesterday vs today] Table 2 compares the sign posted at each beach with the water sampled on the posting day, for 18,998 pairs of consecutive days. The rows give the sign, “Warning” or “Safe”, which is based on the previous day's samples; the columns give whether the posting day's water was over or under Toronto's limit. The sign was right on 1,002 days with a “Warning” sign and water over the limit, and on 14,584 days with a “Safe” sign and water under it. It was wrong on 1,712 days with a “Warning” sign and water under the limit, and on 1,700 days with a “Safe” sign and water over it. Table 3 gives percentages derived from these counts, as answers to questions relevant to a Toronto beach swimmer.

Table 2. Sign posted at each beach, based on the previous day's samples, against the water on the posting day, for pairs of consecutive days, 2007 to 2026. Each cell gives the number of pairs and its share of all pairs.

Sign posted	Water over the limit on the	
	posting day	Water under the limit
“Warning” sign	1,002 (5.3%)	1,712 (9.0%)
“Safe” sign	1,700 (8.9%)	14,584 (76.8%)

Table 3. What the sign says about the water on the posting day, 2007 to 2026, derived from the counts in Table 2. The sign is based on the previous day’s samples.

Question	Answer
On any posting day, how often was the water over the limit?	14%
When the “Warning” sign was up, how often was the water over the limit?	37%
When the “Warning” sign was up, how often was the water under the limit?	63%
When the “Safe” sign was up, how often was the water over the limit?	10%
When the water was over the limit, how often was the “Warning” sign up?	37%

4 Discussion

[Summary] Placeholder: brief summary of the data and results.

[Why beaches differ] Placeholder: Etobicoke Creek at Marie Curtis, the Humber River near Sunnyside, storm-water outfalls, island against mainland beaches; analyses of the data that could test these. Hypotheses, not tested here.

[Improving warnings] Placeholder: the current testing misses most bad days and posts unnecessary warnings most of the time; a prediction model based on weather; qPCR for same-day testing.

[Weaknesses] Placeholder: days over the limit are computed, not recorded postings; the floor and ceiling understate the gaps; no weather data; Sunnyside’s four sites add noise (effect under one percentage point in simulation).

[Next steps] Placeholder: compare days over the limit with the warnings actually posted: the City’s “Beach Condition Data for Researchers” download (JSON) may record them.

Appendix

A Data cleaning

[Cleaning details] Placeholder: downloaded 22 September 2026; 102,534 rows from 3 June 2007 to 16 September 2026, 98,991 with a result; every removal with counts and reasons (1 implausible reading, 220 results from sites 60W and GP6, 400 results outside the season from Victoria Day to Labour Day) (the reading of 6,191,768; sites 60W and GP6 with a footnote on the City's reply and the download date of 22 September 2026; dates outside the season; days with fewer than four results). Tests run on the cleaned data.

B Checking the method on simulated data

[Simulation check] Placeholder: simulated data with known beach differences, the floor and ceiling; the analysis recovers the ordering; censoring leaves about 65% of a true gap; the site-count effect.

C Pairwise comparisons

[Pairwise table] Placeholder: the 45 pairs with Holm-corrected p-values, from `other/results/beach-pairwise-differences.csv`.

References

1. Sanchez J, Tustin J, Heasley C, Patel M, Kelly J, Habjan A, et al. Region-Specific Associations between Environmental Factors and *Escherichia coli* in Freshwater Beaches in Toronto and Niagara Region, Canada. *International Journal of Environmental Research and Public Health*. 2021;18(23):12841. doi:[10.3390/ijerph182312841](https://doi.org/10.3390/ijerph182312841)
2. City of Toronto. About Beach Water Quality [Internet]. 2025 [cited 2026 Sep 22]. Available from: <https://www.toronto.ca/community-people/health-wellness-care/health-inspections-monitoring/swimsafe/beach-water-quality/about-beach-water-quality/>
3. Toronto Public Health. Toronto Beaches Water Quality [Internet]. Open Data Toronto; 2026 [cited 2026 Sep 22]. Available from: <https://open.toronto.ca/dataset/toronto-beaches-water-quality/>
4. City of Toronto. Beach Water Quality [Internet]. 2026 [cited 2026 Sep 22]. Available from: <https://www.toronto.ca/community-people/health-wellness-care/health-inspections-monitoring/swimsafe/beach-water-quality/>
5. City of Toronto. Open Data Toronto [Internet]. 2026 [cited 2026 Sep 22]. Available from: <https://open.toronto.ca/>
6. Saleem F, Schellhorn HE, Simhon A, Edge TA. Same-day *Enterococcus* qPCR results of recreational water quality at two Toronto beaches provide added public health protection and reduced beach days lost. *Canadian Journal of Public Health*. 2023;114:676–87. doi:[10.17269/s41997-023-00763-8](https://doi.org/10.17269/s41997-023-00763-8)
7. Python Software Foundation. Python: A Programming Language [Internet]. 2026. Available from: <https://www.python.org/>
8. Vink R et al. Polars: Blazingly Fast DataFrames [Internet]. 2026. Available from: <https://pola.rs/>
9. Harris CR, Millman KJ, Walt SJ van der, Gommers R, Virtanen P, Cournapeau D, et al. Array programming with NumPy. *Nature*. 2020;585(7825):357–62. doi:[10.1038/s41586-020-2649-2](https://doi.org/10.1038/s41586-020-2649-2)
10. Seabold S, Perktold J. statsmodels: Econometric and Statistical Modeling with Python. In: *Proceedings of the 9th python in science conference*. 2010. p. 92–6. doi:[10.25080/Majora-92bf1922-011](https://doi.org/10.25080/Majora-92bf1922-011)

11. Iannone R et al. pointblank: Data Validation in Python [Internet]. 2026. Available from: <https://posit-dev.github.io/pointblank/>
12. Hunter JD. Matplotlib: A 2D Graphics Environment. Computing in Science & Engineering. 2007;9(3):90–5. doi:[10.1109/MCSE.2007.55](https://doi.org/10.1109/MCSE.2007.55)
13. Arribas-Bel D et al. contextily: Context Geo-Tiles in Python [Internet]. 2026. Available from: <https://contextily.readthedocs.io/>
14. Allaire JJ, Teague C, Scheidegger C, Xie Y, Dervieux C. Quarto [Internet]. 2026. Available from: <https://quarto.org/>
15. Public Health Ontario. Public Beach Water [Internet]. 2024 [cited 2026 Sep 22]. Available from: <https://www.publichealthontario.ca/en/Laboratory-Services/Test-Information-Index/Public-Beach-Water>
16. Young I, Sanchez JJ, Desta BN, Heasley C, Tustin J. Recreational water exposures and illness outcomes at a freshwater beach in Toronto, Canada: A prospective cohort pilot study. PLOS ONE. 2023;18(6):e0286584. doi:[10.1371/journal.pone.0286584](https://doi.org/10.1371/journal.pone.0286584)