

E. coli* at Toronto's Beaches: Are Some Beaches Less Safe Than Others?

Marie Curtis Park East and Sunnyside exceed the swimming limit on a third of summer days, more than seven times as often as Gibraltar Point

Sean Murphy

September 26, 2026

Sentence 1: the general area (Toronto's beaches and the daily *E. coli* warning). Sentence 2: data and method (98,370 samples at ten beaches, 2007 to 2026, from Open Data Toronto; days over the limit, a same-day model, yesterday against today). Sentence 3: headline result (Marie Curtis Park East and Sunnyside over the limit on about a third of days, more than seven times Gibraltar Point; yesterday's samples catch 37% of days over the limit). Sentence 4: implication (knowing which beach you are at tells a swimmer more than yesterday's reading).

1 Introduction

[Why *E. coli*] *E. coli* is a sign of faecal pollution. It lives in the gut of people and other warm-blooded animals, so finding it in lake water means that sewage, storm-water runoff or bird droppings have arrived recently, bringing the viruses and bacteria that make swimmers ill [1]. Swimming in water above Toronto's standard raises the risk of ear, eye, nose, throat and skin infections, and swallowed water can cause stomach illness. Young children, older people and people with weakened immune systems are most at risk [2].

[Questions and gap] Placeholder: the two questions (do some beaches go over the limit more often; how well does yesterday's result predict today's) and the gap: no published comparison of the beaches over 20 years, and no check of the one-day delay.

[Findings] Placeholder: 34% and 32% of days at Marie Curtis Park East and Sunnyside against 4% at Gibraltar Point; a warning based on yesterday catches 37% of the days over the limit.

*Code and data are available at: <https://github.com/seandata8/toronto-beaches-ecoli>.

[Why it matters] Placeholder: which beach matters more than which day, and the posting is a weak guide.

[Roadmap] Placeholder: Section 2 describes the data, how a sample becomes a number, and the methods; Section 3 gives the results; Section 4 what they mean and their limits. The appendices cover cleaning and the check on simulated data.

2 Data

2.1 Source and context

[Data source] To reduce water-borne illness among swimmers, Toronto Public Health, Toronto Water, the police Marine Unit and the City’s parks division jointly sample the water at the city’s ten supervised beaches every day from Victoria Day to Labour Day [3]. Samples are taken at four to six fixed sites along each beach (Figure 1), and when the laboratory count of *E. coli* at a beach is too high, Toronto Public Health posts signs warning against swimming [4]. The results are published on the City’s open data portal [3,5] and form the data for this paper.

[Swimming limit] Toronto’s standard is 100 *E. coli* per 100 mL of water; Ontario and Health Canada use 200 [1,2]. This paper uses 100, the level at which Toronto posts warnings.

[Ethics and alternatives] The data describe water, not people, and contain no personal information. The stakes lie in the advice built on them: a needless warning costs swimmers a beach day, and a missed one exposes them to illness [6]. The City also offers daily swimming conditions for download [4], but this paper uses the open data because they give every result, so each day’s status can be worked out the same way for all twenty years.

[Software] The analysis uses Python [7] with Polars [8], NumPy [9], statsmodels [10], pointblank [11], matplotlib [12] and contextily [13], and the paper is written in Quarto [14].

2.2 Measurement

[Columns] A number in the dataset is the end of a chain that starts with a person wading into the lake. At each site, staff fill a sterile bottle to 200 mL, plunging it to about 30 cm below the surface where people swim. The provincial protocol asks for at least five samples per beach, spread across the swimming area [15]. Most Toronto beaches have five sites; Kew Balmly has six and Sunnyside four. The bottles are kept cold and dark, and must reach a Public Health Ontario laboratory and be tested within a day of collection [15].

[Columns] The laboratory uses membrane filtration [15]. The water is passed through a filter with pores small enough to trap bacteria, and the filter is laid on a growth medium on which *E. coli* colonies can be told apart from other bacteria. Overnight, each *E. coli* cell grows into




Figure 1. Each of Toronto's ten swimming beaches is sampled at four to six sites along the shore. Blue points are the 50 sites this paper uses. The two orange points, sites 60W and GP6, were added in 2026 and sit 14 to 16 km from the beaches they are listed under, so they are left out. Map tiles from Esri.

a colony large enough to see; the colonies are counted and the count is scaled to *E. coli* per 100 mL. The method detects more than 90% of *E. coli* strains, but cells stressed by the lake may not grow and other bacteria can interfere, so a single result carries error [15]. The process takes about a day [2,4], so a result reaches the public a day after the water it describes. When no sample can be taken, usually because of bad weather, the row is kept with its result left blank [2]; 3.5% of rows are blank.

[Floor and ceiling] The way results are reported leaves its mark on the numbers, as Figure 2 shows. Nearly all results (99.9%) are multiples of 10, and half (49%) are exactly 10. Ten is the lowest value the laboratory reports, so a result of 10 means “10 or fewer”: the data have a floor. From 2018 a second limit appears at the top. Results of exactly 1,000 become common while results above 1,000 become rare, and in 2026 none is above 1,000, although before 2018 readings reached 15,820. From 2018, then, a result of 1,000 is best read as “1,000 or more”. Both limits come from the reporting, not the water. The floor hides differences among the clean beaches, and the ceiling hides how bad the worst days were.

[Floor and ceiling] The ceiling matters more once results are averaged. Since 2018, 34 beach-days include a result of exactly 1,000 yet have a geometric mean below 100. On about 8 of them a true reading of 2,000 or less would have put the day over the limit, so the count of days over the limit from 2018 on is, if anything, slightly low.

[Cleaning] The data set was downloaded on 22 September 2026. It has 102,534 rows, one for each site on each sampling day from 3 June 2007 to 16 September 2026, and each row gives the beach, a site code, the date the sample was collected, the *E. coli* result and the site’s location. After leaving out rows with no result, one implausible reading, two sites listed under the wrong beach and dates outside the season (Appendix A), 98,370 results remain, from 50 sites on 2,019 days across 20 years.

[Summary statistics] Table 1 summarises the results at each beach. Results span more than three orders of magnitude, so the figures use a logarithmic scale, on which each equal step is a tenfold increase: 10, 100, 1,000. Most results are low. The median, the middle result when all are put in order, is 20, and at eight of the ten beaches it is 10 or 20. Marie Curtis Park East and Sunnyside stand apart. Their medians are 50 and 60, only about a quarter and a fifth of their results are at the floor, and about a third are above 100. At Gibraltar Point, the cleanest beach, 61% of results are at the floor and 6% are above 100.




Figure 2. Every result the City has published, 2007 to 2026. One point is one water sample (98,991 in total), before any cleaning. Points are spread sideways within each year so that they can be seen; the *E. coli* count itself is never moved. Most results are reported in multiples of 10, which is why the low readings fall into separate rows. The laboratory reports no results as zero, and “no colonies detected” is assumed to be reported as 10. From 2018 on, many results are reported as exactly 1,000, and these are assumed to be readings of 1,000 or more. The dashed line marks 100 *E. coli* per 100 mL, the level at which the City posts a swimming warning. The ringed point is the only reading the analysis removes.

Table 1. Results at each beach, 2007 to 2026. A result is the number of *E. coli* per 100 mL in one water sample. Ten is the lowest value the laboratory reports, and 100 is Toronto’s limit for swimming. The 90th percentile is the level that one result in ten exceeds. Beach-days are days with at least four results. Beaches run from the largest share of results above 100 to the smallest.

Beach	Sites	Results	Me- dian	90th per- centile	At 10	Above 100	Beach- days
Marie Curtis Park East Beach	5	9,612	50	590	27%	35%	1,925
Sunnyside Beach	4	8,006	60	400	20%	34%	1,986
Centre Island Beach	5	9,645	20	220	40%	19%	1,929
Kew Balmy Beach	6	12,047	20	250	47%	19%	2,011
Bluffer’s Beach Park	5	10,006	10	150	55%	13%	2,001
Woodbine Beaches	5	10,035	10	110	59%	11%	2,005
Ward’s Island Beach	5	9,757	10	110	58%	10%	1,952
Cherry Beach	5	10,012	10	90	58%	9%	2,003
Hanlan’s Point Beach	5	9,689	10	80	61%	8%	1,939
Gibraltar Point Beach	5	9,561	10	60	61%	6%	1,911

2.3 Methods

[Beach-days] Toronto’s limit applies not to a single result but to the geometric mean of a beach’s results. A geometric mean of n results multiplies them together and takes the n th root; put another way, it averages the logarithms of the results and converts that average back. It suits counts that span orders of magnitude because a single extreme result does not dominate it. Five results of 10, 10, 10, 10 and 1,000 have an ordinary average of 208 but a geometric mean of 25.

[Beach-days] In this paper a beach-day is over the limit when the geometric mean of that day’s results is above 100. Only days with at least four results count, four rather than five so that Sunnyside, with four sites, is included. That gives 19,662 beach-days, of which 2,816 (14%) are over the limit. Figure 3 shows every one. On 22% of beach-days every result was at the floor of 10.

[Beach-days] A day over the limit is not the same as a day the beach was posted, for three reasons. The data do not record whether a warning went up. A study of the program describes posting decisions based on a geometric mean over two days rather than one [6]. And because results take a day, a warning posted today reflects yesterday’s water. “Over the limit” here means that the water sampled that day exceeded 100, judged after the fact.

[Not in outline] Counting days rather than results changes the shares only a little. At every beach the share of days over the limit is 1.2 to 5.5 percentage points below the share of results

above 100: Marie Curtis Park East, for example, was over the limit on 34% of days and had 35% of results above 100. Averaging five results might be expected to pull most days back under the limit, but it does not, because a beach's sites tend to rise and fall together. When one is high, the others usually are too.

[Model] Placeholder: same-day comparison between beaches (log10 daily geometric mean on beach and date, each symbol explained in words); days in the same month treated as related; what a 95% interval means in plain words; Holm's correction for comparing 45 pairs. The floor and ceiling shrink the gaps to about 65% of their true size in simulation (Appendix B).

[Consecutive days] Placeholder: pairs are two beach-days at the same beach on consecutive calendar days (18,998 pairs); what the confusion matrix shows.

3 Results

[Beach comparison] Placeholder: Figure 3; shares from 34.1% and 31.8% down to 4.4%; the same top two in 19 of 20 years (exception 2021); splitting at 2018 barely moves the shares.

[Model estimates] Placeholder: Figure 4; Marie Curtis and Sunnyside 2.2 times the average beach on the same day, not different from each other, different from every other beach; 31 or 32 of 45 pairs differ; the reporting limits make the gaps look smaller than they are. Centre Island ranks above Kew Balmy in the model but below it in share of days, and why.

[Yesterday vs today] Placeholder: Table 2 (two-by-two from `other/results/warning-confusion-matrix.csv`); a warning based on yesterday catches 37% of days over the limit, and 37% of such warnings are needed; 14% of days are over the limit overall, 37% after a warning, 10% after an all-clear (about 3.5 times as likely). Leave out overall accuracy.

Table 2. Yesterday's advice against today's water at the same beach, 2007 to 2026 (18,998 pairs of consecutive days). A warning means yesterday's geometric mean was above Toronto's limit of 100 *E. coli* per 100 mL; an all-clear means it was not. Each cell gives the number of pairs and its share of all pairs.

Yesterday's advice	Today over the limit	Today under the limit
Warning	1,002 (5.3%)	1,712 (9.0%)
All-clear	1,700 (8.9%)	14,584 (76.8%)

4 Discussion

[Summary] Placeholder: brief summary of the data and results.




Figure 3. How often each beach went over the limit, 2007 to 2026. Each bar counts the beach-days whose geometric mean falls in that range (19,662 days in total; days with fewer than four results are left out). Orange bars, right of the dashed line, are days above Toronto's limit of 100 *E. coli* per 100 mL. The lowest bin includes the 17 days whose mean falls below the reporting floor of 10. Panels run from the largest share of days over the limit to the smallest.




Figure 4. Placeholder caption: each beach’s typical level against the average beach on the same day, on the log10 scale, with 95% intervals. The gaps are understated: the reporting floor and ceiling shrink them to about two thirds of their true size.

[Why beaches differ] Placeholder: Etobicoke Creek at Marie Curtis, the Humber River near Sunnyside, storm-water outfalls, island against mainland beaches; analyses of the data that could test these. Hypotheses, not tested here.

[Improving warnings] Placeholder: the current testing misses most bad days and posts unnecessary warnings most of the time; a prediction model based on weather; qPCR for same-day testing.

[Weaknesses] Placeholder: days over the limit are computed, not recorded postings; the floor and ceiling understate the gaps; no weather data; Sunnyside’s four sites add noise (effect under one percentage point in simulation).

[Next steps] Placeholder: compare days over the limit with the warnings actually posted: the City’s “Beach Condition Data for Researchers” download (JSON) may record them.

Appendix

A Data cleaning

[Cleaning details] Placeholder: every removal with counts and reasons (the reading of 6,191,768; sites 60W and GP6 with a footnote on the City's reply and the download date of 22 September 2026; dates outside the season; days with fewer than four results). Tests run on the cleaned data.

B Checking the method on simulated data

[Simulation check] Placeholder: simulated data with known beach differences, the floor and ceiling; the analysis recovers the ordering; censoring leaves about 65% of a true gap; the site-count effect.

C Pairwise comparisons

[Pairwise table] Placeholder: the 45 pairs with Holm-corrected p-values, and the check with partial months merged. Table from `other/results/beach-pairwise-differences.csv` and `beach-pairwise-grouping-check.csv`.

References

1. Sanchez J, Tustin J, Heasley C, Patel M, Kelly J, Habjan A, et al. Region-Specific Associations between Environmental Factors and *Escherichia coli* in Freshwater Beaches in Toronto and Niagara Region, Canada. *International Journal of Environmental Research and Public Health*. 2021;18(23):12841. doi:[10.3390/ijerph182312841](https://doi.org/10.3390/ijerph182312841)
2. City of Toronto. About Beach Water Quality [Internet]. 2025 [cited 2026 Sep 22]. Available from: <https://www.toronto.ca/community-people/health-wellness-care/health-inspections-monitoring/swimsafe/beach-water-quality/about-beach-water-quality/>
3. Toronto Public Health. Toronto Beaches Water Quality [Internet]. Open Data Toronto; 2026 [cited 2026 Sep 22]. Available from: <https://open.toronto.ca/dataset/toronto-beaches-water-quality/>
4. City of Toronto. Beach Water Quality [Internet]. 2026 [cited 2026 Sep 22]. Available from: <https://www.toronto.ca/community-people/health-wellness-care/health-inspections-monitoring/swimsafe/beach-water-quality/>
5. City of Toronto. Open Data Toronto [Internet]. 2026 [cited 2026 Sep 22]. Available from: <https://open.toronto.ca/>
6. Saleem F, Schellhorn HE, Simhon A, Edge TA. Same-day *Enterococcus* qPCR results of recreational water quality at two Toronto beaches provide added public health protection and reduced beach days lost. *Canadian Journal of Public Health*. 2023;114:676–87. doi:[10.17269/s41997-023-00763-8](https://doi.org/10.17269/s41997-023-00763-8)
7. Python Software Foundation. Python: A Programming Language [Internet]. 2026. Available from: <https://www.python.org/>
8. Vink R et al. Polars: Blazingly Fast DataFrames [Internet]. 2026. Available from: <https://pola.rs/>
9. Harris CR, Millman KJ, Walt SJ van der, Gommers R, Virtanen P, Cournapeau D, et al. Array programming with NumPy. *Nature*. 2020;585(7825):357–62. doi:[10.1038/s41586-020-2649-2](https://doi.org/10.1038/s41586-020-2649-2)
10. Seabold S, Perktold J. statsmodels: Econometric and Statistical Modeling with Python. In: *Proceedings of the 9th python in science conference*. 2010. p. 92–6. doi:[10.25080/Majora-92bf1922-011](https://doi.org/10.25080/Majora-92bf1922-011)

11. Iannone R et al. pointblank: Data Validation in Python [Internet]. 2026. Available from: <https://posit-dev.github.io/pointblank/>
12. Hunter JD. Matplotlib: A 2D Graphics Environment. Computing in Science & Engineering. 2007;9(3):90–5. doi:[10.1109/MCSE.2007.55](https://doi.org/10.1109/MCSE.2007.55)
13. Arribas-Bel D et al. contextily: Context Geo-Tiles in Python [Internet]. 2026. Available from: <https://contextily.readthedocs.io/>
14. Allaire JJ, Teague C, Scheidegger C, Xie Y, Dervieux C. Quarto [Internet]. 2026. Available from: <https://quarto.org/>
15. Public Health Ontario. Public Beach Water [Internet]. 2024 [cited 2026 Sep 22]. Available from: <https://www.publichealthontario.ca/en/Laboratory-Services/Test-Information-Index/Public-Beach-Water>